

LS2-Scalaradd:

(a)

- When the loop is unrolled, there still exist sequential dependencies (because the result of every addition is being put in the same variable!).
- To overcome these dependencies, we increase the number of accumulators. And, as expected the performance is better.
- This is because; the more number of independent instructions available, the pipeline of the corresponding execution unit can be filled (in this case, the pipeline of the floating point unit!).
- The following results are based on running the code on a compute node on Trestles.
- Each node has 4 sockets; each socket has 8 core 2.4GHz AMD – 6136 [K10] Magny-Cours processors.
- I checked the assembly file generated by the compiler (gcc) and found that the instructions being used for the floating point add was ADDSD.

Latency = 4cycles; cycles/issue = 1.

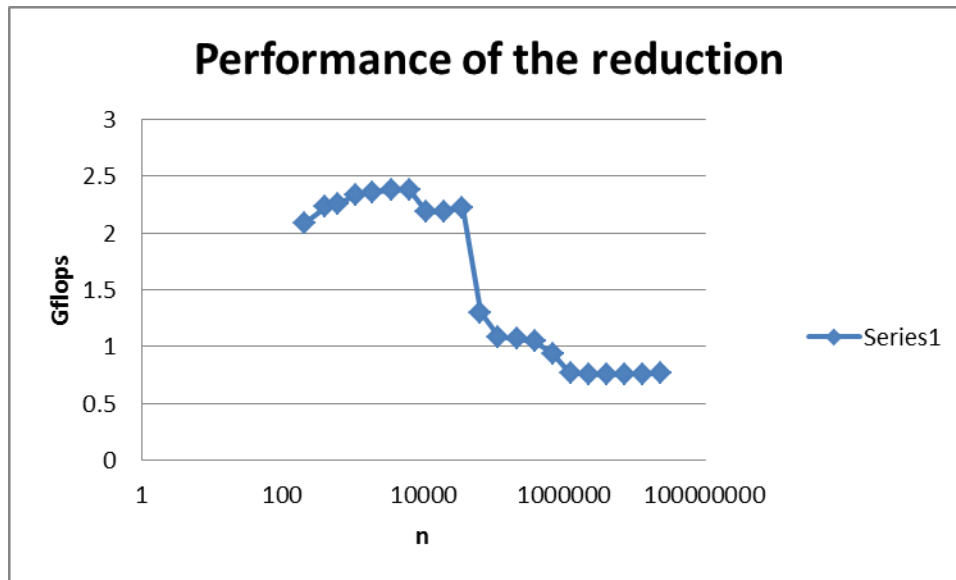
[Source: http://www.agner.org/optimize/instruction_tables.pdf]

Accumulators/Unrolling factor	1	2	3	4	5	6
1	4.017137	4.020056	4.017092	4.020903	4.021047	4.017824
2	0	2.016405	0	2.01686	0	2.016865
3	0	0	1.349423	0	0	1.351381
4	0	0	0	1.016129	0	0
5	0	0	0	0	1.017868	0
6	0	0	0	0	0	1.019907

Maximal Performance for Accumulators (K) = 4, Unrolling (L) = 4

K = 4 is enough to fill the pipeline of the execution unit.

(b)



- The plateaus in the above graph correspond to different performance levels of the processor.
- As the size of the vector (n) increases, it can be observed that the performance reduces.
- Hence, this reduction in performance can be attributed to the performance of the memory hierarchy of the processor.
- The **first plateau**, where the performance is almost equal to the theoretical peak [i.e. $1\text{flop/cycle} * 2.4\text{GHz} = 2.4 \text{ Gflops}$], for $i = 1$ through 7 which means $n = 200$ to 6200 corresponds to the case where all the elements of the vector can fit in the **L1 cache** !!
- The **second plateau** after the 'small' drop corresponds to the case when the vector size spills the data to the **L2 cache**. $n = 11100, 19900, 35800$.
- After the big drop, the **third plateau** corresponds to the case when the data spills to the **L3 cache**. $n = 64300$ to 674700
- The **fourth plateau** corresponds to the case when there are **DRAM accesses**. $n = 674700$ to 22946900.
- The big drop after the plateau corresponding to L2 cache can be attributed to **greater number of misses in the TLB** with increase in size of n. [virtual address translation]

- **Note:** I could have given figures of number of L1, L2, L3 cache misses, but there is no 'perf' on Trestles and I was taking time to get started with 'papi' !!